



TCP & UDP



2.5 TCP and UDP

- + **How does this support my pentesting career?**
 - TCP Session Attacks
 - Advanced DoS attacks
 - Network scanning

2.5 TCP and UDP

- + In this section, you will see how the **transport layer** works, and how the application layer uses its services to identify server and client processes.
- + The **Transmission Control Protocol (TCP)** and the **User Datagram Protocol (UDP)** are the most common transport protocols used on the Internet.

2.5 TCP and UDP

- + Before checking out the different services that a transport layer protocol can offer to the application layer, let's consider something important about networks.
- + Computer networks can be **unreliable**. This means that some **packets can be lost** during their trip from source to destination. A packet can be lost because of network congestion, temporary loss of connection and other technical issues.

2.5 TCP and UDP

- + When designing a transport layer protocol, the designer must choose how to deal with these limitations. For example, TCP:
 - **Guarantees packet delivery.** Because of that, an application that needs a guaranteed delivery will use TCP as the transport protocol.
 - Is also **connection oriented**. It must establish a connection before transferring data.
- + Keep in mind these facts during your study!

2.5 TCP and UDP

- + TCP is the most used transport protocol on the Internet. The vast majority of applications use it, and the IP protocol suite is often called **TCP/IP**.
- + Email clients, web browsers and FTP clients are some common applications using TCP.

2.5 TCP and UDP

- + On the other hand, UDP is much more simple than TCP:
 - + It does **not guarantee** packet delivery.
 - + It is **connectionless**.

2.5 TCP and UDP

- + UDP is faster than TCP, as it provides a **better throughput** (number of packets per second); in fact, it is used by **multimedia applications** that can tolerate packet loss but are throughput intensive.
- + For example, UDP is used for VoIP and video streaming: applications where you can tolerate a little glitch in the audio or video.

2.5 TCP and UDP

- + Here we can see a comparison table between TCP and UDP.

TCP	UDP
Lower throughput	Better throughput
Connection-oriented	Connectionless
Guarantees delivery	Does not guarantee packet delivery

2.5.1 Ports


- + Applications and their processes use TCP and UDP to send and receive data over the network. When an IP datagram reaches a host, how can the transport layer **know what the destination process is?**
- + We'll now introduce **ports**.

2.5.1 Ports

- + **Ports** are used to identify a single network **process** on a machine. If you want to unequivocally identify a process on a network, you need to know the **<IP> : <Port> pair**.
- + As an example, you can compare the port to the recipient's name on a letter; the street address (IP) identifies the building, while the person name identifies the final recipient of the letter.


2.5.1.1 Ports Examples

- + In this image, you can see how every client application on *Client PC* uses a different port.




2.5.1.1 Ports Examples

- + The browser uses local port 3028 to connect to the web server...



2.5.1.1 Ports Examples

+ ... while the mail client uses local port 1022.



2.5.1.1 Ports Examples

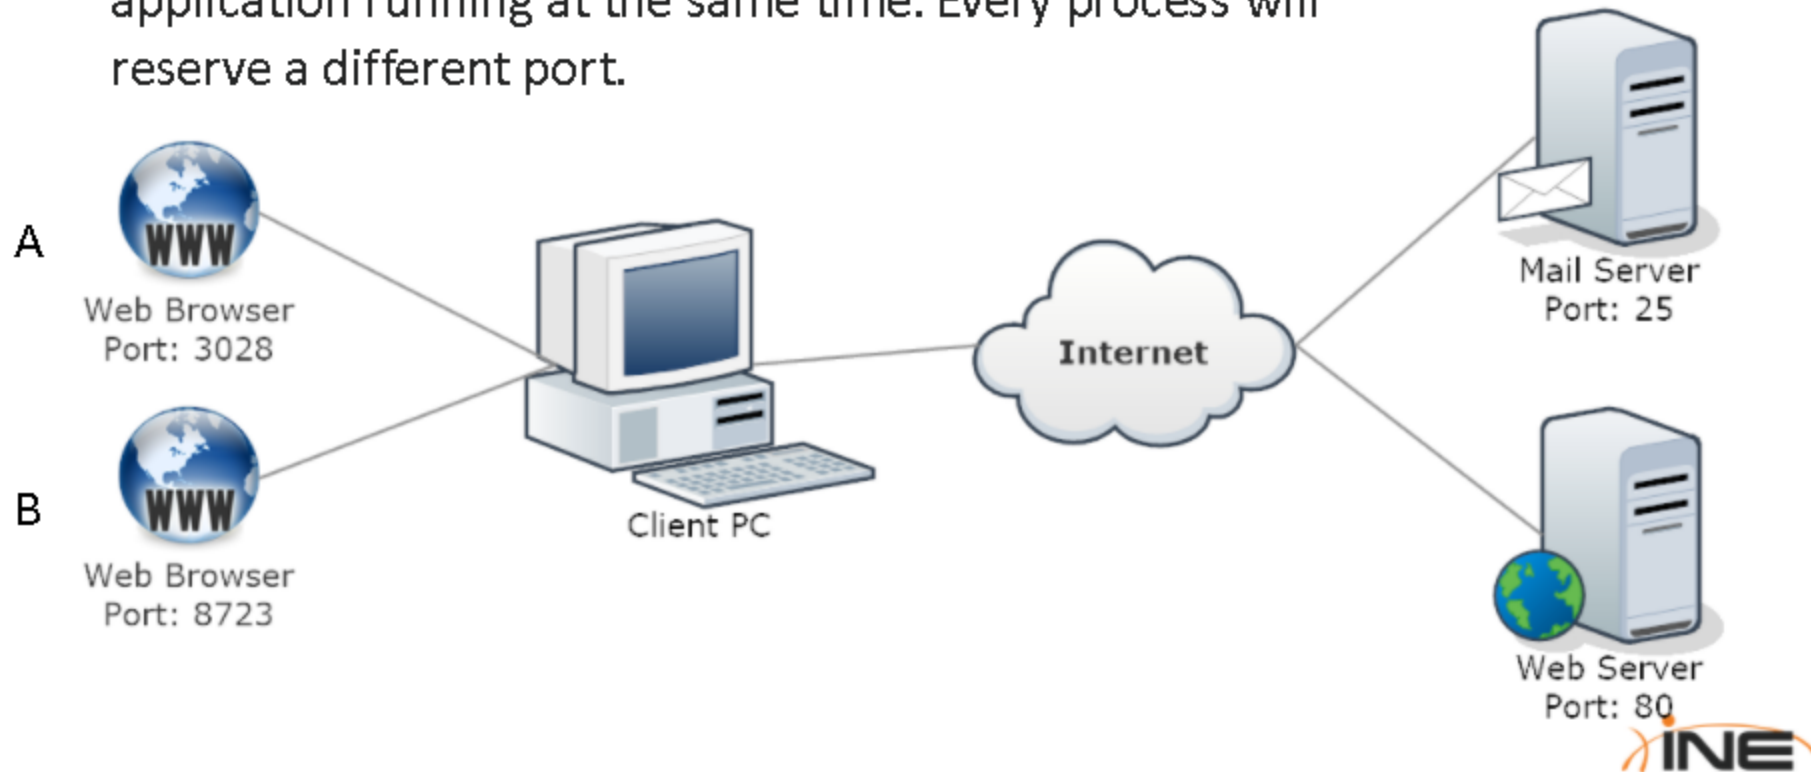
- + In the previous example:
 - All the communication from the web browser to the web server will have 3028 as the source port and 80 as the destination port.
 - All the communication back from the web server to the browser will have 80 as the source port and 3028 as the destination port.

2.5.1.1 Ports Examples

- + Similarly, for the mail client and server:
 - All the communication from the mail client to the server will have 1022 as the source port and 25 as the destination port.
 - All the communication back from the mail server to the mail client will have 25 as the source port and 1022 as the destination port.


2.5.1.1 Ports Examples

- Furthermore, you may also have multiple instances of the same application running at the same time. Every process will reserve a different port.




2.5.1.1 Ports Examples

- + In this example, 'A' communicates with the web server using 3028 as the source port...



2.5.1.1 Ports Examples

+ ... while 'B' uses port 8723.



2.5.1 Ports

- + To correctly address a process on a network, you have to refer to the `<IP>:<Port>` pair. For example:
 - + `192.168.5.3:80`
 - + `10.11.12.1:443`
 - + `172.16.8.9:22`
- + But, how can you know the right port for a common service?

2.5.2 Well-known Ports

- + Ports in the ranging from **0-1023**, the first 1024 that is, are called **well-known ports** and are used by servers for the most common services.
- + For example, when a web browser connects to a server via HTTPS, the user does not have to manually specify 443 as the destination port.

2.5.2 Well-known Ports

- + Each common protocol has a well-known port in the 0-1023 range. Common server processes, or daemons, use well-known ports most of the time.
- + Ports are assigned by IANA and are referenced in [this document](http://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml).

2.5.2 Well-known Ports

+ You do not need to know all the service port assignments, but you should at least remember the most common, such as:

- SMTP (25)
- SSH (22)
- POP3 (110)
- IMAP (143)
- HTTP (80)
- HTTPS (443)
- NETBIOS (137, 138, 139)
- SFTP (115)
- Telnet (23)
- FTP (21)
- RDP (3389)
- MySQL (3306)
- MS SQL Server (1433)

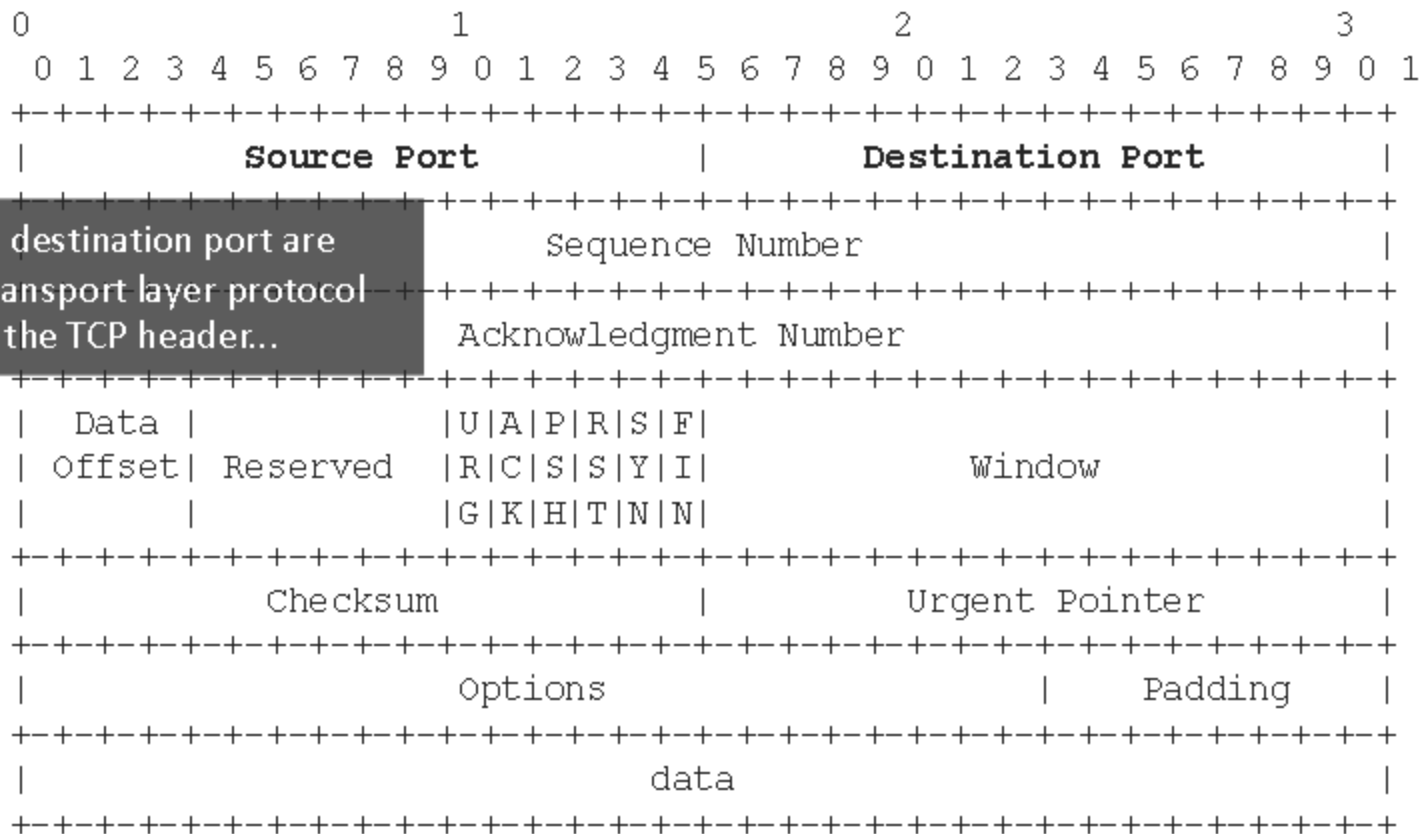
2.5.2 Well-known Ports

- + As briefly introduced before, a **daemon** is a program that runs a service. System administrators can change the daemon configuration, **changing the port** the service listens to for connection. They do that to make services recognition a little bit harder for hackers.
- + For example, you could find an FTP daemon listening on port 4982 instead of 21 or SSH listening on port 8821.

2.5.3 TCP and UDP headers

- + Let's now see how ports are used by applications.
- + How can server and client applications know which port to use?
They use two fields in the TCP or UDP header: the **source** and **destination** ports.

2.5.3.1 TCP Header



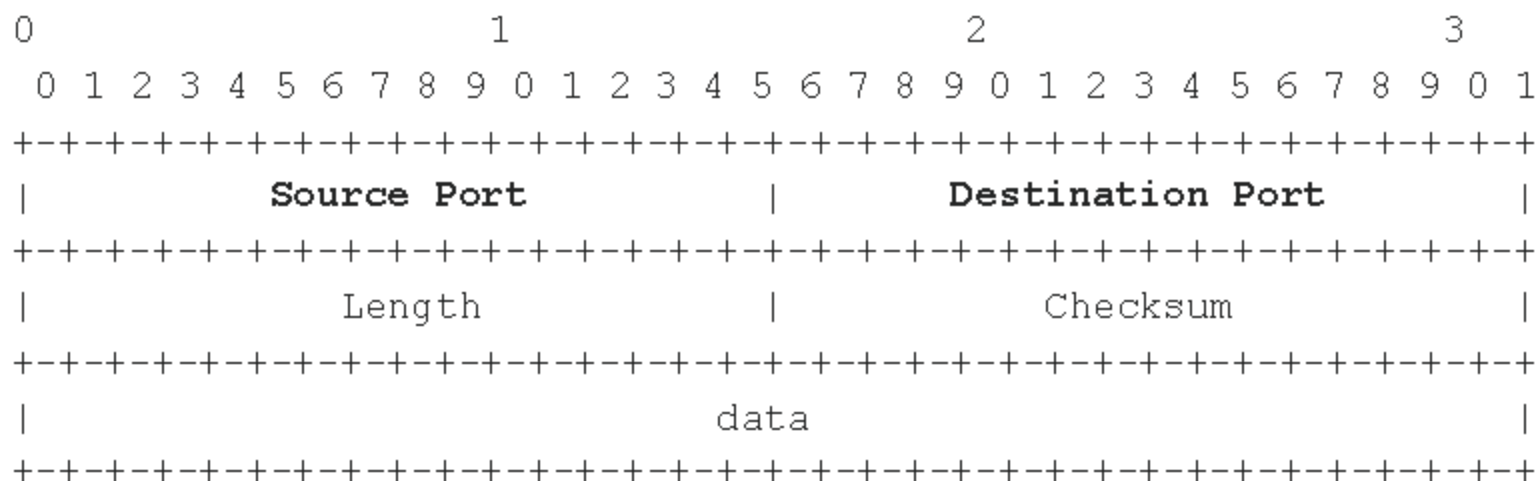
The source and destination port are included in the transport layer protocol header. Like the TCP header..

2.5.3.1 TCP Header


[illegible]

2.5.3.2 UDP Header

...or the UDP header.



2.5.3.2 UDP Header



2.5.4 Netstat Command

- + To check the listening ports and the current (TCP) connections on a host you can use:
 - + `netstat -ano` on Windows
 - + `netstat -tunp` on Linux
 - + `netstat -p tcp -p udp` together with `lsof -n -i4TCP -i4UDP` on MacOS
- + Use these commands to show information about the processes listening on the machine and processes connecting to remote servers.

2.5.4 Netstat Command

- + Another great tool for Windows is [TCPView](#) from Sysinternals.
- + TCPView shows:
 - Process name
 - PID
 - Protocol
 - Local and remote addresses
 - Local and remote ports
 - State of the connection (if applicable)

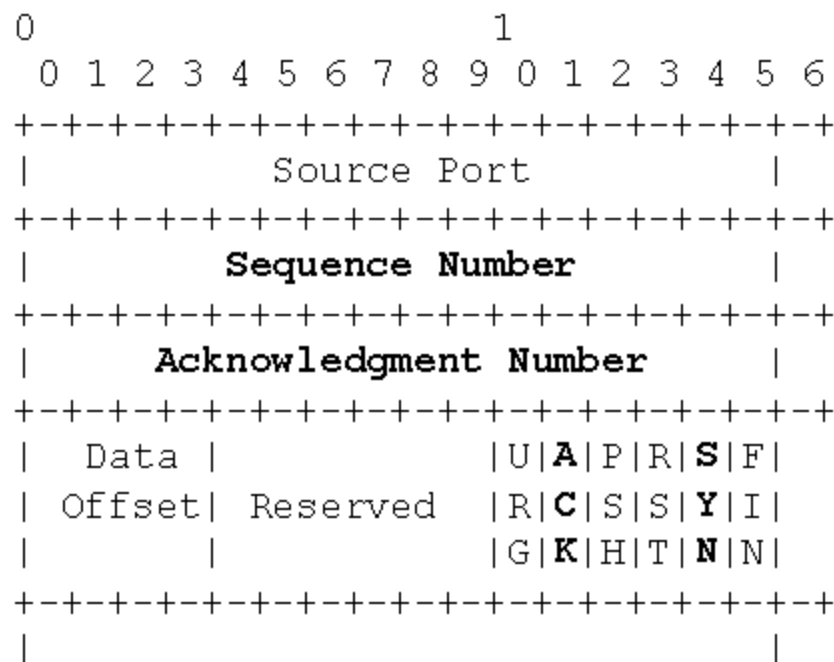
2.5.5 TCP Three Way Handshake

- + We have seen that TCP is **connection oriented**. Now, let's look at how TCP connections work, as well as highlight the most important factors involved, from the penetration tester's point of view, in a 3-way handshake.
- + To establish a connection between two hosts running TCP, they must perform three steps: the **three-way handshake**. They can then start the actual data transmission.

2.5.5 TCP Three Way Handshake


The header fields involved
in the handshake are:

- + Sequence number
- + Acknowledgement numbers
- + SYN and ACK flags




2.5.5 TCP Three Way Handshake

- + The steps in the handshake are used to synchronize the sequence and acknowledgment numbers between the server and the client.




2.5.5 TCP Three Way Handshake

- + During the first step, the client sends a TCP packet to the server with the SYN flag enabled and a random sequence number.




2.5.5 TCP Three Way Handshake

- + In the second step, the server replies by sending a packet with both the SYN and ACK flag set and another random sequence number.



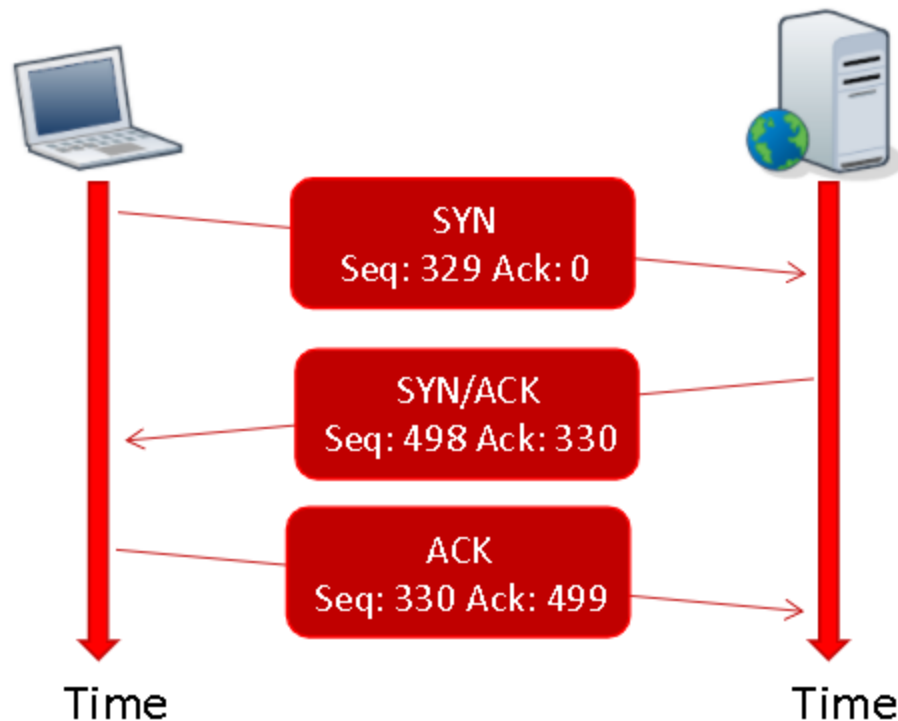
2.5.5 TCP Three Way Handshake

- + The ACK number is always a simple increment of the SYN number sent by the client.



2.5.5 TCP Three Way Handshake

- + Finally, the client completes the synchronization by sending an ACK packet.
- + Note that the client behaves just like the server when sending ACK packets.



References

- + For additional information, please check out these references:
 - IP Layer Network Administration with Linux.
 - TCP/IP Tutorial and Technical Overview.
 - Packet Analysis Reference Guide v3.0.
- + [Service Name and Transport Protocol Port Number Registry:](http://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml)
<http://www.iana.org/assignments/service-names-port-numbers/service-names-port-numbers.xhtml>
- + [TCPView:](http://technet.microsoft.com/en-us/sysinternals/bb897437) <http://technet.microsoft.com/en-us/sysinternals/bb897437>